

2. Linear Regression with One Variable – or Univariate Linear Regression

⇒ MODEL REPRESENTATION:

Training set of housing prices (Portland, OR)	Size in feet ² (x)	Price (\$) in 1000's (y)
	→ 2104	460
	1416	232
	→ 1534	315
	852	178

Notation:

- m = Number of training examples
- x 's = "input" variable / features
- y 's = "output" variable / "target" variable
- (x, y) - one training example
- $(x^{(i)}, y^{(i)})$ - i^{th} training example

$x^{(1)} = 2104$
 $x^{(2)} = 1416$
 $y^{(1)} = 460$

$m = 47$

Andre

➤ Other Notations:

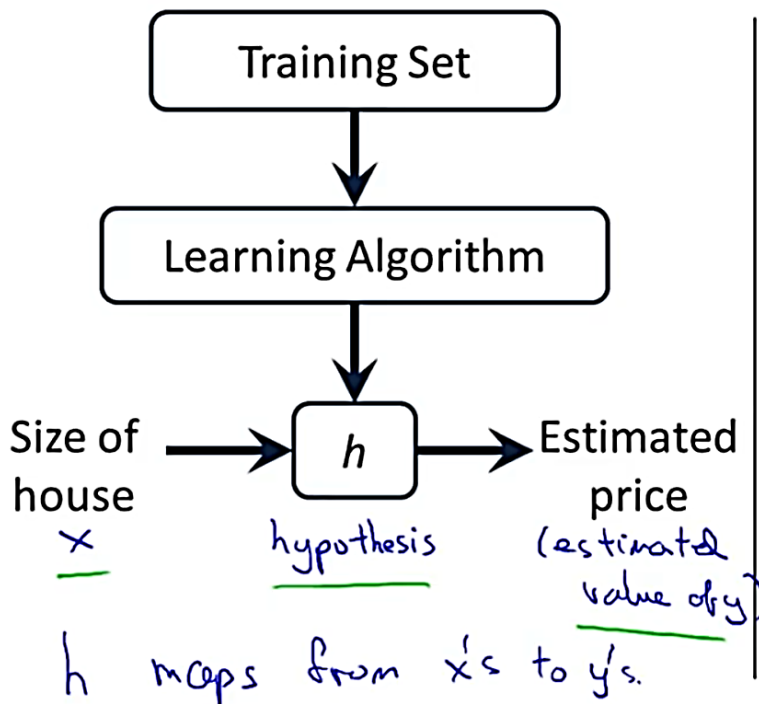
X = space of input values

Y = space of output values

Dataset = list of m training examples → $(x^{(i)}, y^{(i)})$; $i = 1, 2, \dots, m$

Hypothesis Function: Function which is derived by feeding training data (Input and Output (Supervised)) to the learning algorithm, which can then be used to predict o/p for new input data.

- For a supervised problem: $h : X \rightarrow Y$ so that $h(x)$ is a “good” predictor for the corresponding value of y .




08:00 / 08:09

How do we represent h ?

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$


Shorthand: $h(x)$



Linear regression with one variable. (x)
Univariate linear regression.
one variable

COST FUNCTION: Can measure the **accuracy** of our hypothesis function by **Choosing parameters** of $h(x)$ such that **$h(x)$ is close to y** for data in the training set.

Cost function is a **minimization** function:



minimize θ_0, θ_1

$\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

#training examples

$h_{\theta}(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Idea: Choose θ_0, θ_1 so that $h_{\theta}(x)$ is close to y for our training examples (x, y)

x, y

minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1
Cost function
Squared error function

Here, we try to **minimize** the $(1/2m) \times$ (sum of squared differences b/w predicted value and actual value given in dataset).

- “J” is the cost function or **squared error cost function** or **Mean squared error**
- **(1/2m)** is for averaging the squared difference.

Minimize means, we try to find values of **θ** parameters such that the cost function is minimized.

Cost ➔ An **average difference** of all the results of the hypothesis with inputs from x's and the actual output y's.

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i)^2$$

To break it apart, it is $\frac{1}{2} \bar{x}$ where \bar{x} is the mean of the squares of $h_{\theta}(x_i) - y_i$, or the difference between the predicted value and the actual value.

- The mean is **halved** as a convenience for the computation of the **gradient descent**, as the derivative term of the square function will cancel out the 1/2 term.

COST FUNCTION INTUITION: Training set data is scattered on x-y plane. We try to draw a straight line through it. **Goal** ➔ **find best fitting line.**

Ideally, the line should pass through all the points of our training data set. In such a case, the value of “J” will be 0.

- For simplicity: let $h_{\theta}(x) = \theta_1 x$

➔ $h_{\theta}(x)$
(for fixed θ_1 , this is a function of x)

➔ $J(\theta_1)$
(function of the parameter θ_1)


For $\Theta=1 \rightarrow h(\Theta)=x \rightarrow J(\Theta)=0$

For $\Theta=0.5 \rightarrow h(\Theta)=0.5x \rightarrow$


$J(\Theta)=0.58$

For $\Theta=0 \rightarrow h(\Theta)=0 \rightarrow J(\Theta)=2.3$

- $J(\Theta)$ is the average of sq of diff bw $h(x)$ and y :
- $h(x)$ = **predicted** value at given training data
- y = **actual** value of o/p in training data




- **Vertical lines** represent the given **difference**



- For different values of Θ , we try to minimize **$J(\Theta)$ {error}**, which occurs at $\Theta=1$. Therefore, we choose $\Theta=1$ as our best fitting curve: $h(x)=x$


>> For more **complex** $h(x)$ fcn like **$h(x) = \Theta_0 + \Theta_1 x$** : we have to plot **$J(\Theta_0, \Theta_1)$** in **3D**. As, for different combinations of Θ_0 and Θ_1 , J can be different.



These can be more easily represented using **contour figures**: A contour plot is a graph that contains many contour lines. A **contour line** of a two variable function has a **constant value** at all points on the line.

$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




Graph b/w θ_0 and θ_1 : these **ellipses** are the combinations of θ_0 and θ_1 for which value of J is same. Points other than on ellipses are also valid points, they also correspond to a unique value of J .

- The **best** combination (one which **minimizes** $J(\Theta_0, \Theta_1)$) of Θ_0 and Θ_1 lies around **center** of the innermost circle.


⇒ **GRADIENT DESCENT**: an algorithm to minimize $J(\Theta_0, \Theta_1)$

- Start with some Θ_0, Θ_1
- Keep changing Θ_0, Θ_1 to reduce J until minimum is reached



Here, we start at a value of Θ_0, Θ_1 and keep going down on $J(\Theta_0, \Theta_1)$ curve until we reach a **local minima**. We will know that we have succeeded when our cost function is at the very bottom of the pits in our graph.

- If we start at a diff value of Θ_0, Θ_1 , we end having different minima.



We are not graphing x and y itself, but the parameter range of our hypothesis function and the cost resulting from selecting a particular set of parameters.

Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$ (for $j = 0$ and $j = 1$)
}

The **slope of the tangent** is the **derivative** at that point and it will give us a **direction** to move towards. We make steps down the cost function in the direction with the steepest descent. The size of each step is determined by the parameter α , which is called the **learning rate**.

α = Learning rate

SIMULTANEOUS UPDATE: first we calculate new value for both θ_0 and θ_1 , then only we update their values. So order of execution of statements is:

Correct: Simultaneous update

```
temp0 :=  $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$   
temp1 :=  $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$   
 $\theta_0 :=$  temp0  
 $\theta_1 :=$  temp1
```



Incorrect:

```
→ temp0 :=  $\theta_0 - \alpha \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$   
→  $\theta_0 :=$  temp0  
→ temp1 :=  $\theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$  ←  
→  $\theta_1 :=$  temp1
```



Here, if update θ_0 before actually calculating the value of new θ_1 , the θ_0 used in equation of θ_1 will be new θ_0 , not the one we wanted to minimize for.


At each iteration j , one should simultaneously update the parameters $\dots \theta_0, \theta_1 \dots \theta_n$. Updating a specific parameter prior to calculating another one on the j^{th} iteration would yield a wrong implementation.

GRADIENT DESCENT INTUITION: for simplicity we only use one parameter:

$$h(x) = \theta_1 \cdot x$$

→ α is positive


- ⇒ For a value of θ_1 , if the **slope** of $J(\theta)$ is **positive**: θ_1 decreases
- ⇒ For **negative slope** of $J(\theta)$: θ_1 increases
- ⇒ θ_1 eventually **converges** to its minimum.



- ⇒ If the value of α is **too small**: gradient descent takes baby steps towards the min.


$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

If α is too small, gradient descent can be slow.



⇒ If α is too large: gradient descent takes huge steps.. in such case the gradient descent may even **overshoot the min** if the diff b/w initial Θ and Θ_{\min} is less than the value of **jump in Θ ($\alpha * \text{derivative of } J$)** and it may start going further and further from the min.

If α is too large, gradient descent can overshoot the minimum. It may fail to converge, or even diverge.




- Therefore, we should adjust our parameter to ensure that the gradient descent algorithm **converges** in a **reasonable time**.
 - If the Θ is already at its **local minimum**, the **slope will be 0**.. thus Θ won't change
- ⇒ Even if the learning rate α is fixed, the **slope gets smaller as we reach** towards the **minima**.. so the steps automatically become smaller

Gradient descent can converge to a local minimum, even with the learning rate α fixed.

$$\theta_1 := \theta_1 - \alpha \frac{d}{d\theta_1} J(\theta_1)$$

As we approach a local minimum, gradient descent will automatically take smaller steps. So, no need to decrease α over time.



⇒ For a **linear regression** model: **Derivative of J:**

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) &= \frac{\partial}{\partial \theta_j} \left[\frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \right] \\ &= \frac{\partial}{\partial \theta_j} \frac{1}{2m} \sum_{i=1}^m (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2 \end{aligned}$$

repeat until convergence: {


$$\begin{aligned} \theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x_i) - y_i) \\ \theta_1 &:= \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m ((h_{\theta}(x_i) - y_i) x_i) \\ &\} \end{aligned}$$

⇒ **For θ_0** – derivate wrt θ_0

⇒ **For θ_1** – derivate wrt θ_1


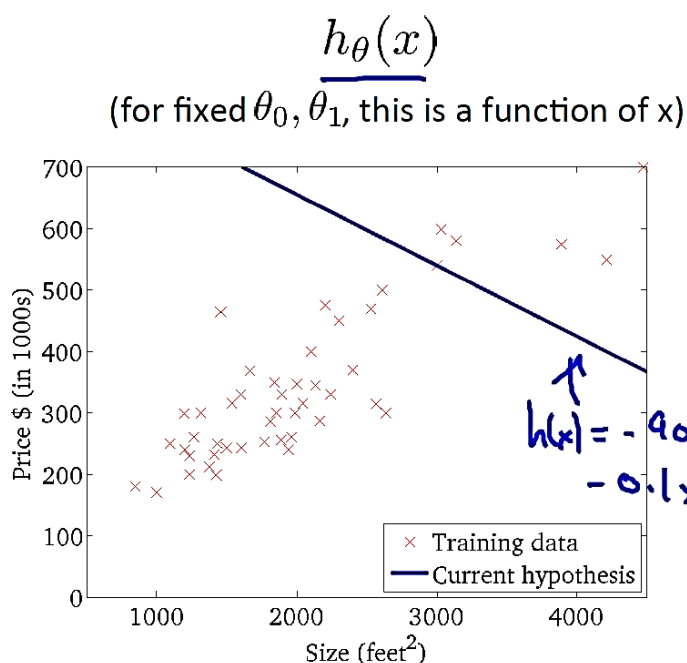
- For a **linear regression** model: the curve is always a **convex curve** (bowl shaped).

It has only **one optimum** → **global minima** (assuming the learning rate α is not too large).



- In the contour curve: we start of with any value of θ_0 and θ_1 and then we minimize the J.
- **We approach the min as we reach towards the center.**


We start at an **arbitrary** value for θ_0 and θ_1 :



We start minimizing $J(\Theta_0, \Theta_1)$ with our gradient descent algo:


$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)



J is a complicated quadratic function

The ellipses shown above are the **contours** of a quadratic function

Batch Gradient Descent → Each step of gradient descent uses all training examples.

The point of all this is that if we start with a guess for our hypothesis and then **repeatedly apply these gradient descent** equations, our **hypothesis** will become more and more **accurate**.
